


Introduction to AI

Support Vector Machines



Rodrigo Cabral

EUR DS4H-LIFE-SPECTRUM

cabral@unice.fr

Outline


1. Hard-margin support vector machines for binary classification
2. Soft-margin support vector machines for binary classification
3. Nonlinear support vector machines
4. Multi-class classification
5. Conclusions

1. Hard-margin support vector machines for binary classification
2. Soft-margin support vector machines for binary classification
3. Nonlinear support vector machines
4. Multi-class classification
5. Conclusions

Linearly separable classes

What happens if the classes are separable?

- ▶ Which line do you choose?



- ▶ Logistic regression suffers from two major issues in this case:
 1. $\kappa \rightarrow +\infty$ to have $J \rightarrow 0$: optimization algorithms become unstable.
 2. The solution of the problem is not unique.


Are there other approaches different from regularized logistic regression?

Large margin linear decisions

What would be a good decision line?

- ▶ Line with a “safe zone” around it.
 - ▶ Correct classification for new data points which can be located around training data.
 - ▶ Good generalization properties.

- ▶ Margin: maximum width around the decision line before hitting a data point.

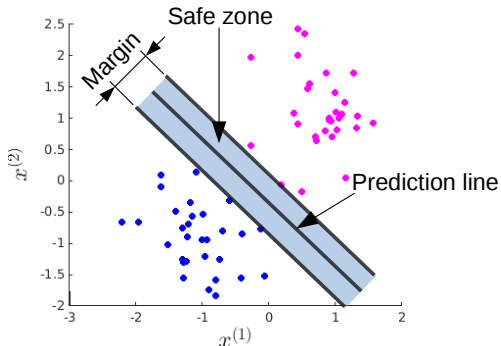


Large margin linear decisions

What would be a good decision line?

- ▶ Line with a “safe zone” around it.
 - ▶ Correct classification for new data points which can be located around training data.
 - ▶ Good generalization properties.


- ▶ Margin: maximum width around the decision line before hitting a data point.
- ▶ Choose decision line with **largest margin**.



Reminder: hyperplane equations and half-spaces

Equation of an hyperplane
passing through 0

$$\mathbf{x}\mathbf{w} = 0$$
$$\|\mathbf{w}\|_2 = 1 \text{ - unitary vector}$$




Reminder: hyperplane equations and half-spaces

Equation of an hyperplane
passing through 0

$$\mathbf{x}\mathbf{w} = 0$$


$\|\mathbf{w}\|_2 = 1$ - unitary vector



General hyperplane equation

$$\mathbf{x}\mathbf{w} + b = 0$$

b - signed distance of closest
point to origin




Reminder: hyperplane equations and half-spaces

Equation of an hyperplane
passing through 0

$$\mathbf{x}\mathbf{w} = 0$$


$\|\mathbf{w}\|_2 = 1$ - unitary vector



General hyperplane equation

$$\mathbf{x}\mathbf{w} + b = 0$$

b - signed distance of closest
point to origin




Upper half-space

$$\mathbf{x}\mathbf{w} + b \geq 0$$


Lower half-space

$$\mathbf{x}\mathbf{w} + b \leq 0$$



Large margin linear decisions

Possible decision line and margin



- **Objective:** choose line parameters w and b leading to maximum margin 2δ such that data points are outside the margin.

Large margin linear decisions

Problem formulation as constrained minimization

- ▶ The margin constraint for a data point \mathbf{x}_i is

$$\begin{cases} \mathbf{x}_i \mathbf{w} + b \geq \delta & \text{if } y_i = 1 \\ \mathbf{x}_i \mathbf{w} + b \leq -\delta & \text{if } y_i = 0 \end{cases}$$

- ▶ Dividing both sides by δ , defining $\beta' = \mathbf{w}/\delta$ and $\beta_0 = b/\delta$, we have

$$\begin{cases} \mathbf{x}_i \beta' + \beta_0 \geq 1 & \text{if } y = 1 \\ \mathbf{x}_i \beta' + \beta_0 \leq -1 & \text{if } y = 0 \end{cases}$$

- ▶ This can be rewritten in a more compact form as follows

$$(2y_i - 1)(\mathbf{x}_i \beta' + \beta_0) \geq 1$$

Large margin linear decisions

Problem formulation as constrained minimization

- Note that $\delta = 1/\|\beta'\|_2$, therefore the optimization problem we want to solve is

$$\text{maximize} \quad \frac{2}{\|\beta'\|_2}$$

$$\text{with respect to} \quad \beta', \beta_0$$

$$\begin{aligned} \text{subject to} \quad & (2y_i - 1)(\mathbf{x}_i\beta' + \beta_0) \geq 1 \\ & \text{for all } i \in \{1, \dots, N\} \end{aligned}$$

Large margin linear decisions

Problem formulation as convex optimization

- ▶ The previous problem can be transformed into the following equivalent problem:

$$\begin{array}{ll}\text{minimize} & \frac{\|\beta'\|_2^2}{2} \\ \text{with respect to} & \beta', \beta_0 \\ \text{subject to} & (2y_i - 1)(\mathbf{x}_i\beta' + \beta_0) \geq 1 \\ & \text{for all } i \in \{1, \dots, N\}\end{array}$$

- ▶ This is a convex optimization problem. More precisely a problem from the class of quadratic programs.
- ▶ It does not have a closed-form solution.
 \implies Fortunately, many numerical optimization algorithms can be used to solve it.

Large margin linear decisions

Dual formulation

- ▶ In practice, this optimization problem is very complex to solve. Moreover, the solution cannot be interpreted.
- ▶ The problem is recast in its **Lagrangian dual form**:

$$\text{maximize} \quad \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j (2y_i - 1)(2y_j - 1) \mathbf{x}_i \mathbf{x}_j^T$$

with respect to

$$\alpha_1, \dots, \alpha_N$$

subject to

$$\alpha_1 \geq 0, \dots, \alpha_N \geq 0$$

$$\sum_{i=1}^N \alpha_i (2y_i - 1) = 0$$

- ▶ α_i are called the Lagrangian dual variables.

Large margin linear decisions

Solution and prediction

- ▶ It can be shown that the optimal solution $\hat{\beta}', \beta_0$ of the initial problem is then written as a function of the solution of the dual α_i :

$$\hat{\beta}' = \sum_{i=1}^N \alpha_i (2y_i - 1) \mathbf{x}_i^T$$

and for any i for which $\alpha_i > 0$ we can retrieve $\hat{\beta}_0$ by solving

$$(2y_i - 1) (\mathbf{x}_i \beta' + \beta_0) = 1$$

- ▶ **Prediction:**

$$\hat{y}(\mathbf{x}) = \begin{cases} 1, & \text{if } f_{\beta}(\mathbf{x}) \geq 0 \\ 0, & \text{if } f_{\beta}(\mathbf{x}) < 0 \end{cases}$$


where $f_{\beta}(\mathbf{x}) = \mathbf{x} \hat{\beta}' + \hat{\beta}_0 = \sum_{i=1}^N \alpha_i (2y_i - 1) \mathbf{x} \mathbf{x}_i^T + \hat{\beta}_0$.

Support vector machine

Support vectors

- ▶ It can be shown that $\alpha_i > 0$, only if \mathbf{x}_i lies exactly on the optimal margin boundary.
- ▶ The prediction line is determined only by these \mathbf{x}_i , which are closer to the border.
 - ⇒ Decision line is defined only by most ambiguous observations.
- ▶ If you remove one of these points from the data set, the decision line may change. These data vectors “support” the decision line.
 - ⇒ They are called **support vectors** and the method is called **support vector machine (SVM)**.


Example of linear SVM classifier




1. Hard-margin support vector machines for binary classification
2. Soft-margin support vector machines for binary classification
3. Nonlinear support vector machines
4. Multi-class classification
5. Conclusions

Classes not fully separable

What do we do if the classes are not linearly separable ?




- ▶ The data is almost linearly separable, except for a few points.
- ▶ We will deal with this case first.



- ▶ The data is not close to linear separability, at least with these features.

Soft margin SVM

Soft version of the problem including “slacks” on the margins



- ▶ We soften the margins by allowing slacks $\xi_i \geq 0$.
- ▶ We should try to keep slacks as small as possible by penalizing them.
 - \implies Still trying to keep decision line as far as possible from the two classes.

Soft margin SVM

Problem formulation as convex optimization

- ▶ We have only slight modifications of the hard margin problem:

$$\text{minimize} \quad \frac{\|\beta'\|_2^2}{2} + C \sum_{i=1}^N \xi_i$$

$$\text{with respect to} \quad \beta', \beta, \xi_1, \dots, \xi_N$$

$$\begin{aligned} \text{subject to} \quad & (2y_i - 1)(\mathbf{x}_i\beta' + \beta_0) \geq 1 - \xi_i \quad i \in \{1, \dots, N\} \\ & \xi_1 \geq 0, \dots, \xi_N \geq 0 \end{aligned}$$

- ▶ Coefficient C allows to control trade-off between margin maximization and fitting to data.

Soft margin SVM

Dual formulation

- It can be shown that the dual formulation only changes slightly:

$$\text{maximize} \quad \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j (2y_i - 1)(2y_j - 1) \mathbf{x}_i \mathbf{x}_j^T$$

with respect to

$$\alpha_1, \dots, \alpha_N$$

subject to

$$0 \leq \alpha_i \leq C \text{ for all } \alpha_i$$

$$\sum_{i=1}^N \alpha_i (2y_i - 1) = 0$$

Soft margin SVM

Dual formulation

- It can be shown that the dual formulation only changes slightly:

$$\text{maximize} \quad \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j (2y_i - 1)(2y_j - 1) \mathbf{x}_i \mathbf{x}_j^T$$

with respect to

$$\alpha_1, \dots, \alpha_N$$

subject to

$$0 \leq \alpha_i \leq C \text{ for all } \alpha_i$$

$$\sum_{i=1}^N \alpha_i (2y_i - 1) = 0$$

- We still have

$$\hat{\beta}' = \sum_{i=1}^N \alpha_i (2y_i - 1) \mathbf{x}_i^T$$

but $\hat{\beta}_0$ is different:



$$\hat{\beta}_0 = (2y_k - 1)(1 - \xi_k) - \mathbf{x}_k \hat{\beta}' \quad \text{with } k = \arg \max_i \alpha_i$$

Soft margin SVM

Support vectors and effect of C

- ▶ Support vectors are still \mathbf{x}_i corresponding to positive α_i .
- ▶ The predictions are given in the same way as before.

Same example but with different C




- ▶ Support vectors are either on the margin border or beyond it.

1. Hard-margin support vector machines for binary classification
2. Soft-margin support vector machines for binary classification
3. Nonlinear support vector machines
4. Multi-class classification
5. Conclusions

Nonlinear separation

What do we do when classes are not directly linearly separable?

- ▶ Example of data set with classes which are not linearly separable in their original features:




- ▶ We can do the same as we did in linear regression to fit nonlinear curves: add transformations of the features.
⇒ The border seems quadratic, we can try adding $(x^{(1)})^2$ to the features.

Nonlinear separation with feature transformation

Transformed feature space

- Transformation of each data feature with a function $\phi(\mathbf{x})$ generates an augmented feature space:

$$\mathbf{x}' = \phi(\mathbf{x}) = \begin{bmatrix} x^{(1)} & x^{(2)} & (x^{(1)})^2 \end{bmatrix}$$




- Classes can be separated by a plane.

⇒ Linear separation!

SVM with transformed observations

What do we get if we apply SVM with this extended feature space?



- ▶ Classes can be perfectly separated with a hard-margin SVM.

SVM with transformed observations

Dual formulation

- For a general transformation $\phi(\mathbf{x}) : \mathbb{R}^M \rightarrow \mathbb{R}^L$, the dual formulation of the soft SVM changes to

$$\text{maximize} \quad \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j (2y_i - 1)(2y_j - 1) \phi(\mathbf{x}_i) \phi^T(\mathbf{x}_j)$$

with respect to

$$\alpha_1, \dots, \alpha_N$$

subject to

$$0 \leq \alpha_i \leq C \text{ for all } \alpha_i$$

$$\sum_{i=1}^N \alpha_i (2y_i - 1) = 0$$

SVM with transformed observations

Dual formulation

- ▶ For a general transformation $\phi(\mathbf{x}) : \mathbb{R}^M \rightarrow \mathbb{R}^L$, the dual formulation of the soft SVM changes to

$$\text{maximize} \quad \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j (2y_i - 1)(2y_j - 1) \phi(\mathbf{x}_i) \phi^T(\mathbf{x}_j)$$

with respect to

$$\alpha_1, \dots, \alpha_N$$

subject to

$$0 \leq \alpha_i \leq C \text{ for all } \alpha_i$$

$$\sum_{i=1}^N \alpha_i (2y_i - 1) = 0$$

- ▶ The function $f_{\beta}(\mathbf{x})$ used for prediction changes to

$$f_{\beta}(\mathbf{x}) = \mathbf{x} \hat{\beta}' + \hat{\beta}_0 = \sum_{i=1}^N \alpha_i (2y_i - 1) \phi(\mathbf{x}) \phi^T(\mathbf{x}_i) + \hat{\beta}_0$$

and $\hat{\beta}_0$ is

$$\hat{\beta}_0 = (2y_k - 1)(1 - \xi_k) - \sum_{i=1}^N \alpha_i (2y_i - 1) \phi(\mathbf{x}_k) \phi^T(\mathbf{x}_i) \quad \text{with } k = \arg \max_i \alpha_i$$

Kernels

Kernels

- ▶ The transformed features appear always through scalar products

$$k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})\phi^{\top}(\mathbf{x}')$$

- ▶ For any chosen $\phi(\cdot)$, $k(\mathbf{x}, \mathbf{x}')$ is a **kernel function**.
- ▶ Kernel functions measure similarity between vectors.

Kernels

Kernels

- ▶ The transformed features appear always through scalar products

$$k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})\phi^T(\mathbf{x}')$$

- ▶ For any chosen $\phi(\cdot)$, $k(\mathbf{x}, \mathbf{x}')$ is a **kernel function**.
- ▶ Kernel functions measure similarity between vectors.
- ▶ What if we do not know the transformation $\phi(\cdot)$?
 \implies Choose directly the kernel function.

Kernels

- ▶ The transformed features appear always through scalar products

$$k(\mathbf{x}, \mathbf{x}') = \phi(\mathbf{x})\phi^T(\mathbf{x}')$$

- ▶ For any chosen $\phi(\cdot)$, $k(\mathbf{x}, \mathbf{x}')$ is a **kernel function**.
- ▶ Kernel functions measure similarity between vectors.
- ▶ What if we do not know the transformation $\phi(\cdot)$?
 \implies Choose directly the kernel function.
- ▶ Examples of kernel functions:
 - ▶ Radial basis function (RBF): $k(\mathbf{x}, \mathbf{x}') = \exp(-\gamma\|\mathbf{x} - \mathbf{x}'\|_2^2)$
 - ▶ Exponential: $k(\mathbf{x}, \mathbf{x}') = \exp(-\gamma\|\mathbf{x} - \mathbf{x}'\|_2)$
 - ▶ d -th degree polynomial: $k(\mathbf{x}, \mathbf{x}') = (\gamma\mathbf{x}\mathbf{x}'^T + r)^d$

Kernel SVM

Dual formulation with kernels - Kernel SVM

- For a kernel $k(\cdot, \cdot)$, the dual formulation of the soft SVM becomes

$$\text{maximize} \quad \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j (2y_i - 1)(2y_j - 1) k(\mathbf{x}_i, \mathbf{x}_j)$$

with respect to

$$\alpha_1, \dots, \alpha_N$$

subject to

$$0 \leq \alpha_i \leq C \text{ for all } \alpha_i$$

$$\sum_{i=1}^N \alpha_i (2y_i - 1) = 0$$

- The function $f_{\beta}(\mathbf{x})$ used for prediction changes to

$$f_{\beta}(\mathbf{x}) = \mathbf{x} \hat{\beta}' + \hat{\beta}_0 = \sum_{i=1}^N \alpha_i (2y_i - 1) k(\mathbf{x}, \mathbf{x}_i) + \hat{\beta}_0$$

and $\hat{\beta}_0$ is


$$\hat{\beta}_0 = (2y_k - 1)(1 - \xi_k) - \sum_{i=1}^N \alpha_i (2y_i - 1) k(\mathbf{x}_k, \mathbf{x}_i) \quad \text{with } k = \arg \max_i \alpha_i$$

Kernel SVM


Kernel value as a function of distance

- Let us define the distance $d = \|\mathbf{x} - \mathbf{x}'\|_2$.

RBF: $k(d) = \exp(-\gamma d^2)$



Exponential: $k(d) = \exp(-\gamma d)$





- $\gamma > 0$ is an hyperparameter indicating how support vectors will influence decisions on their neighborhoods.
 - Large γ produces local influence.
 - Small γ produces global influence.

Kernel SVM

Example

- ▶ Previous nonlinear separation example with hard margin kernel SVM: RBF kernel with two different γ .




- ▶ γ needs to be chosen wisely depending on the complexity of the decision boundary.

1. Hard-margin support vector machines for binary classification
2. Soft-margin support vector machines for binary classification
3. Nonlinear support vector machines
- 4. Multi-class classification**
5. Conclusions

Multi-class classification problem

Problem with K classes

- ▶ Given classes with K different labels $y \in \{0, \dots, K-1\}$, how do we define the decision boundaries to generate the predictions?



- ▶ Three methods: **One vs. one**, **One vs. rest**, **Multi-class**.


Multi-class classification problem

One vs. one (OVO)

1. For each pair of classes (k, k') , $k \neq k'$, learn the binary classifier parameters $\beta_{(k,k')}$ to generate predictions $\hat{y}_{(k,k')}$ disregarding the other classes.
2. Combine binary classifiers via voting mechanism, for example, majority voting:

$$\hat{y}_{\text{OVO}}(\mathbf{x}) = \arg \max_{1 \leq k' \leq K} |\{k : \hat{y}_{(k,k')} = k'\}|$$

where $|\cdot|$ denotes the number of elements of a set.



Drawbacks:


- ▶ Computational: train all combinations of binary classifiers.
- ▶ Overfitting: size of training sample could be small for a given pair.

Multi-class classification problem

One vs. rest (OVR)

1. Learn binary classifier parameters β_k for each class against all the other classes merged.
2. Since $f_{\beta_k}(\mathbf{x})$ is a measure of the depth of \mathbf{x} within class k , prediction is given by

$$\hat{y}_{\text{OVR}}(\mathbf{x}) = \arg \max_{1 \leq k \leq K} f_{\beta_k}(\mathbf{x})$$



Drawback:

- Calibration: classifier functions $f_{\beta_k}(\mathbf{x})$ may not be comparable.

It is however quite simple to implement and less complex than OVO.

Multi-class classification problem

Multi-class approach

- ▶ Both logistic regression and SVM can be modified to explicitly deal with multiple classes.
 - ⇒ No need to learn binary classifiers and to apply fusion rules.
- ▶ We are not going to detail this approach in this class.
 - ⇒ This kind of approach may increase learning complexity, without leading to a significant increase in classification performance with respect to OVO and OVR.

1. Hard-margin support vector machines for binary classification
2. Soft-margin support vector machines for binary classification
3. Nonlinear support vector machines
4. Multi-class classification
5. Conclusions

Conclusions

Logistic regression

- ▶ Logistic regression can be seen as a simple adaptation of linear regression to do classification. Learning is equivalent to solve a smooth optimization problem.
- ▶ If we want nonlinear boundaries we need to include them explicitly.
- ▶ Parameter vector $\hat{\beta}$ allows for interpretation of the importance of the features.
- ▶ Complexity of underlying optimization problem depends on the dimension of feature space.
- ▶ It is sensible to outliers (it is designed for this) and it cannot be used directly for separable problems.

Conclusions

SVM

- ▶ SVM focus directly on the classification problem. It learns a robust separation boundary. Learning is equivalent to solve a quadratic program.
- ▶ The boundary can be turned nonlinear either with feature transformation or with kernel SVM.
 - ⇒ Kernels can be designed for non-numeric data types (graphs, sequences, relational data).
 - ⇒ Kernel SVM has been applied to many different fields ranging from text to genetic data.
- ▶ Dual formulation allows to know which observations are important for the SVM classifier, they are the support vectors.

Conclusions

SVM

- ▶ In the standard SVM approach, $\hat{\beta}'$ allows to analyze the importance of the features.
In kernel SVM, there is no $\hat{\beta}'$, therefore the importance of the features is difficult to be analyzed.
- ▶ Complexity of underlying optimization problem depends on the number of observations mainly.
- ▶ It can be made insensible to outliers (with C) and it can directly deal with separable problems.
- ▶ Tuning kernel SVM parameters C , type of kernel and γ can be quite difficult. Most common approach consists in using cross-validation (topic of more advanced lectures).